



The Kernel

The Kernel

Scheduler: Round-Robin, Priority Queues, Tree Flavours

Scheduler Actors: Features, Timers, Async I/O

Streams Backends: Zero-copy, Message Passing



Linear Backends: Async I/O Disk Streams, Network Streams

Indexed Backends: Timers, Actors

Backpressured Message Bus/Buffers: Arc/Vec prealloc

Class: Low Latency, Real Time

Linear: MQ, EXT, DISK, NET
Trees: TIMERS
Priority Queues: TASKS, IRQ



IO

MIO compatible polling loop based on Readiness Queue

SERVER

POLL

READINESS

NODES

SELECTOR

CONN #1

OS: EPOLL WAIT

CONN #2

EVENTS

EVENT

TOKEN

READY

MQ

Queue Types

SPSC/LINK

4-10ns Lowest Latency Possible

MPSC/SUB

10-40ns Reducer or Subscribe Polling

SPMC/PUB

10-40ns Publisher Multicursor

FAST DELIVERY CASE

Single Threaded Task Configuration
to be compared as reference

I/O TASK



CPU TASK



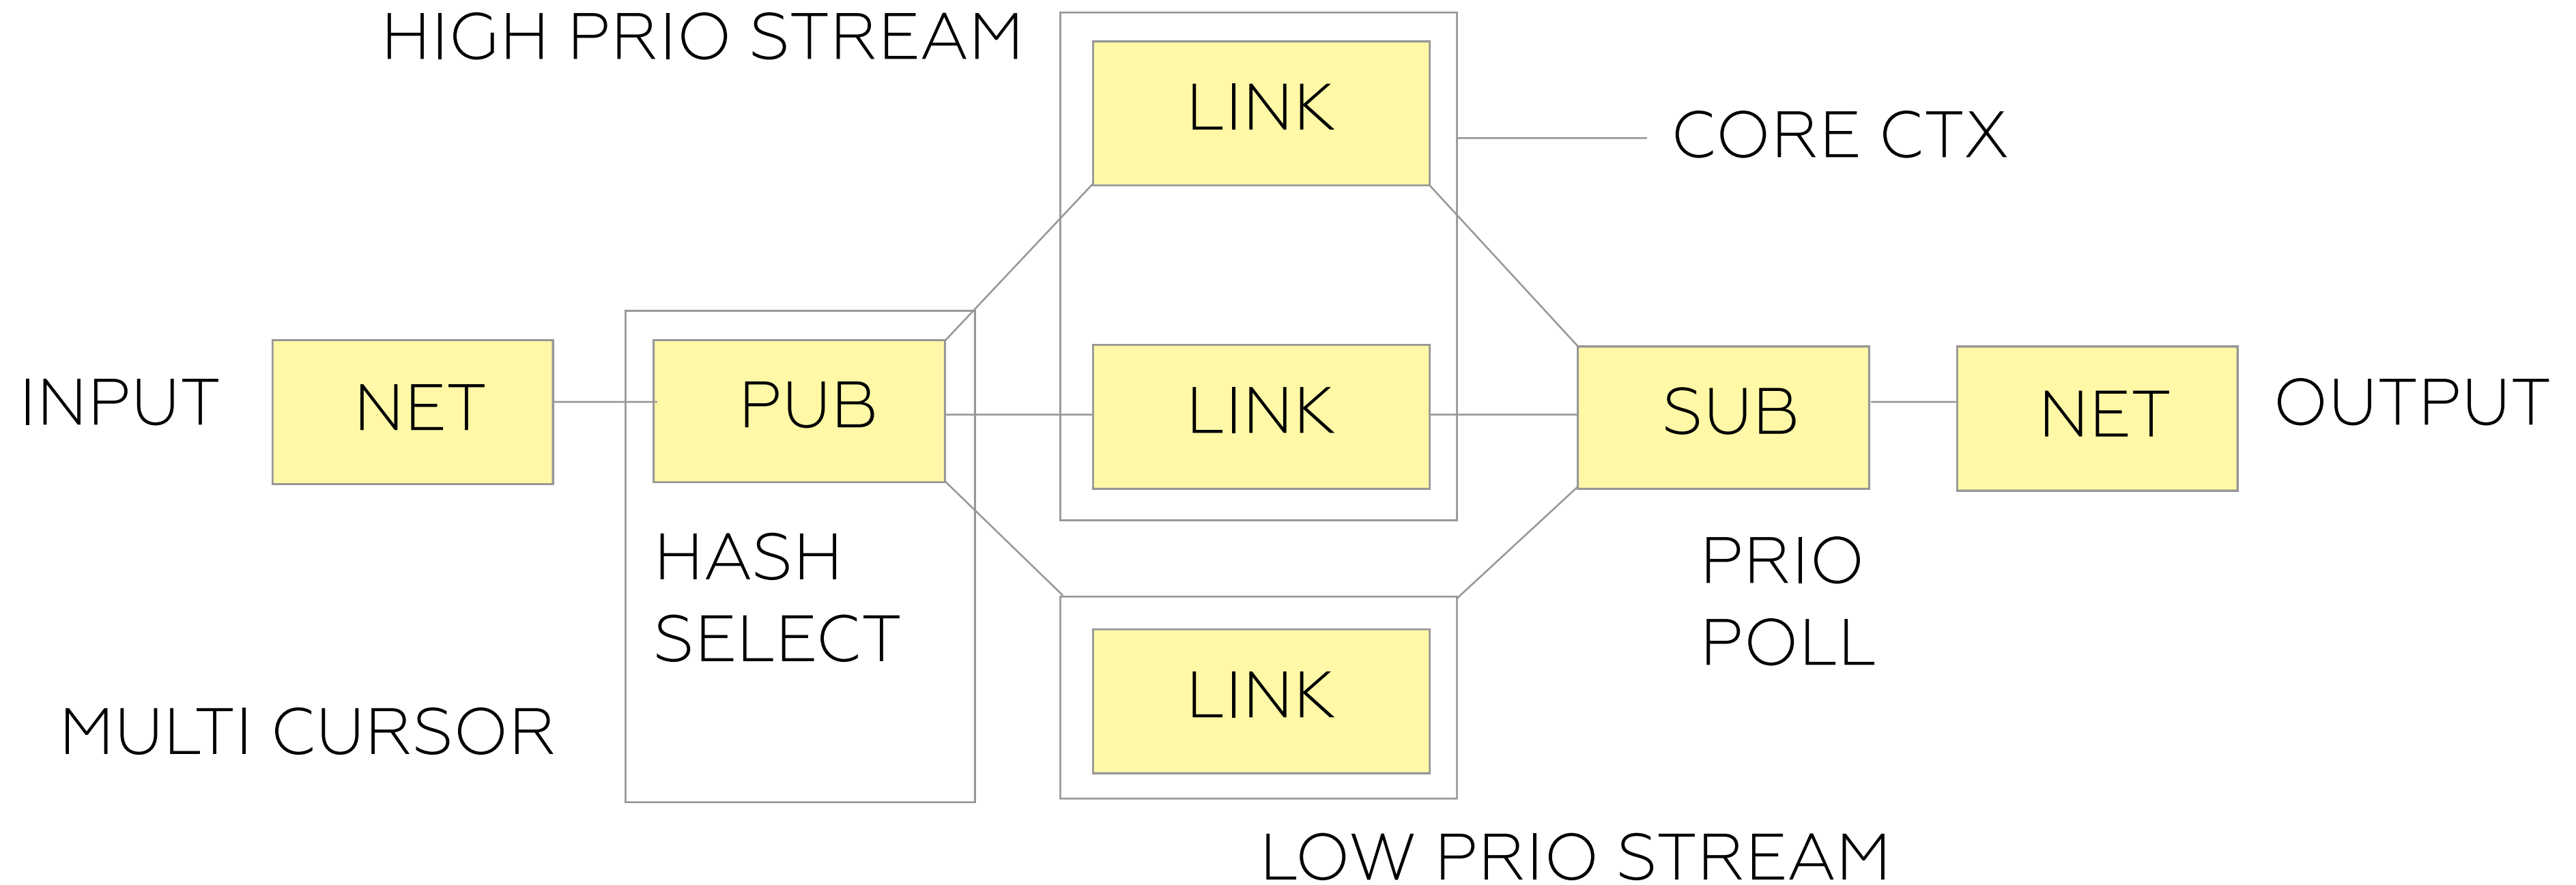
I/O TASK



You can use inplace message modifying and reduce copies to unpack and pack.


LOAD BALANCING CASE

Load Balancing of Priority Streams per Core Buckets



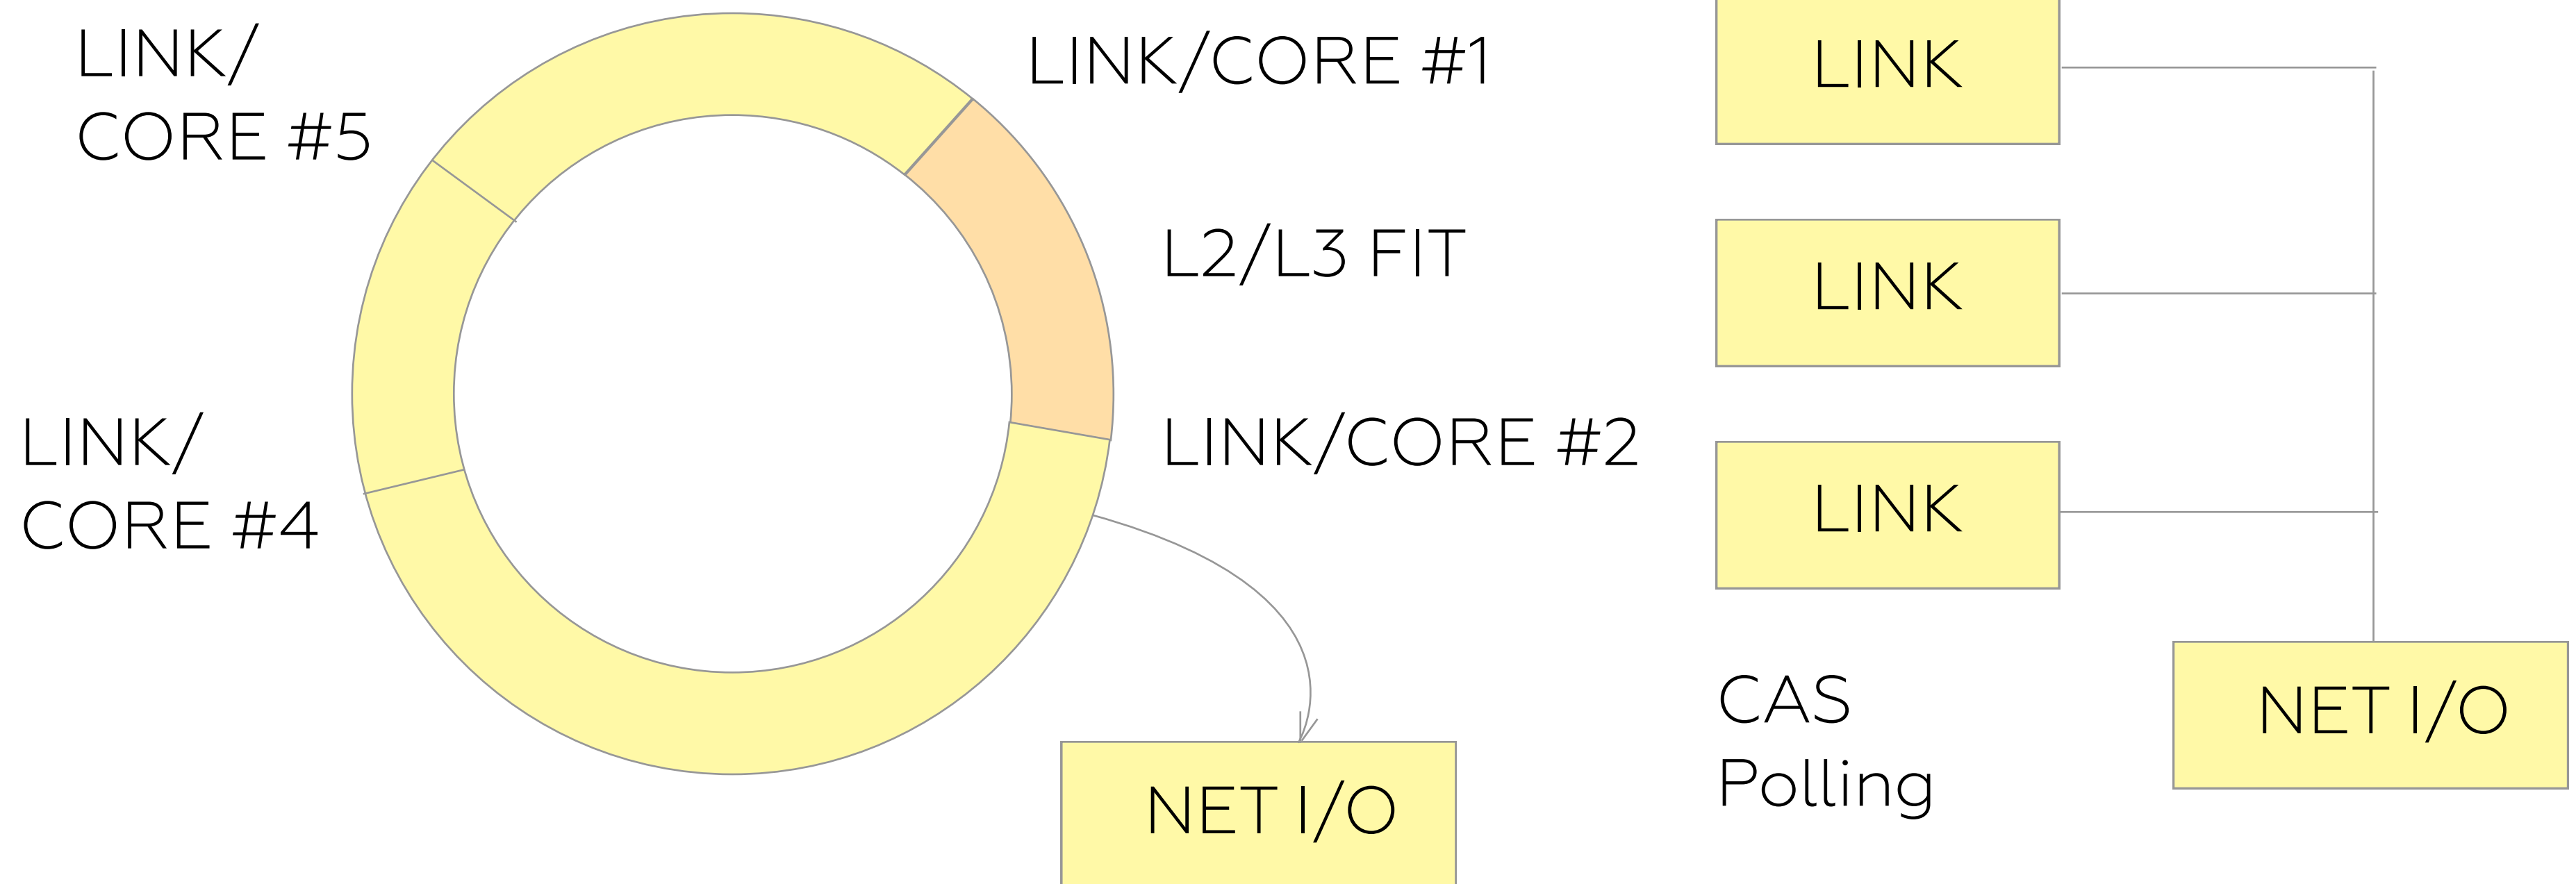
PUBLISHER CASE

PUB Implementation for Zero-Copy Multiple Consumer Publishing (SPMC)



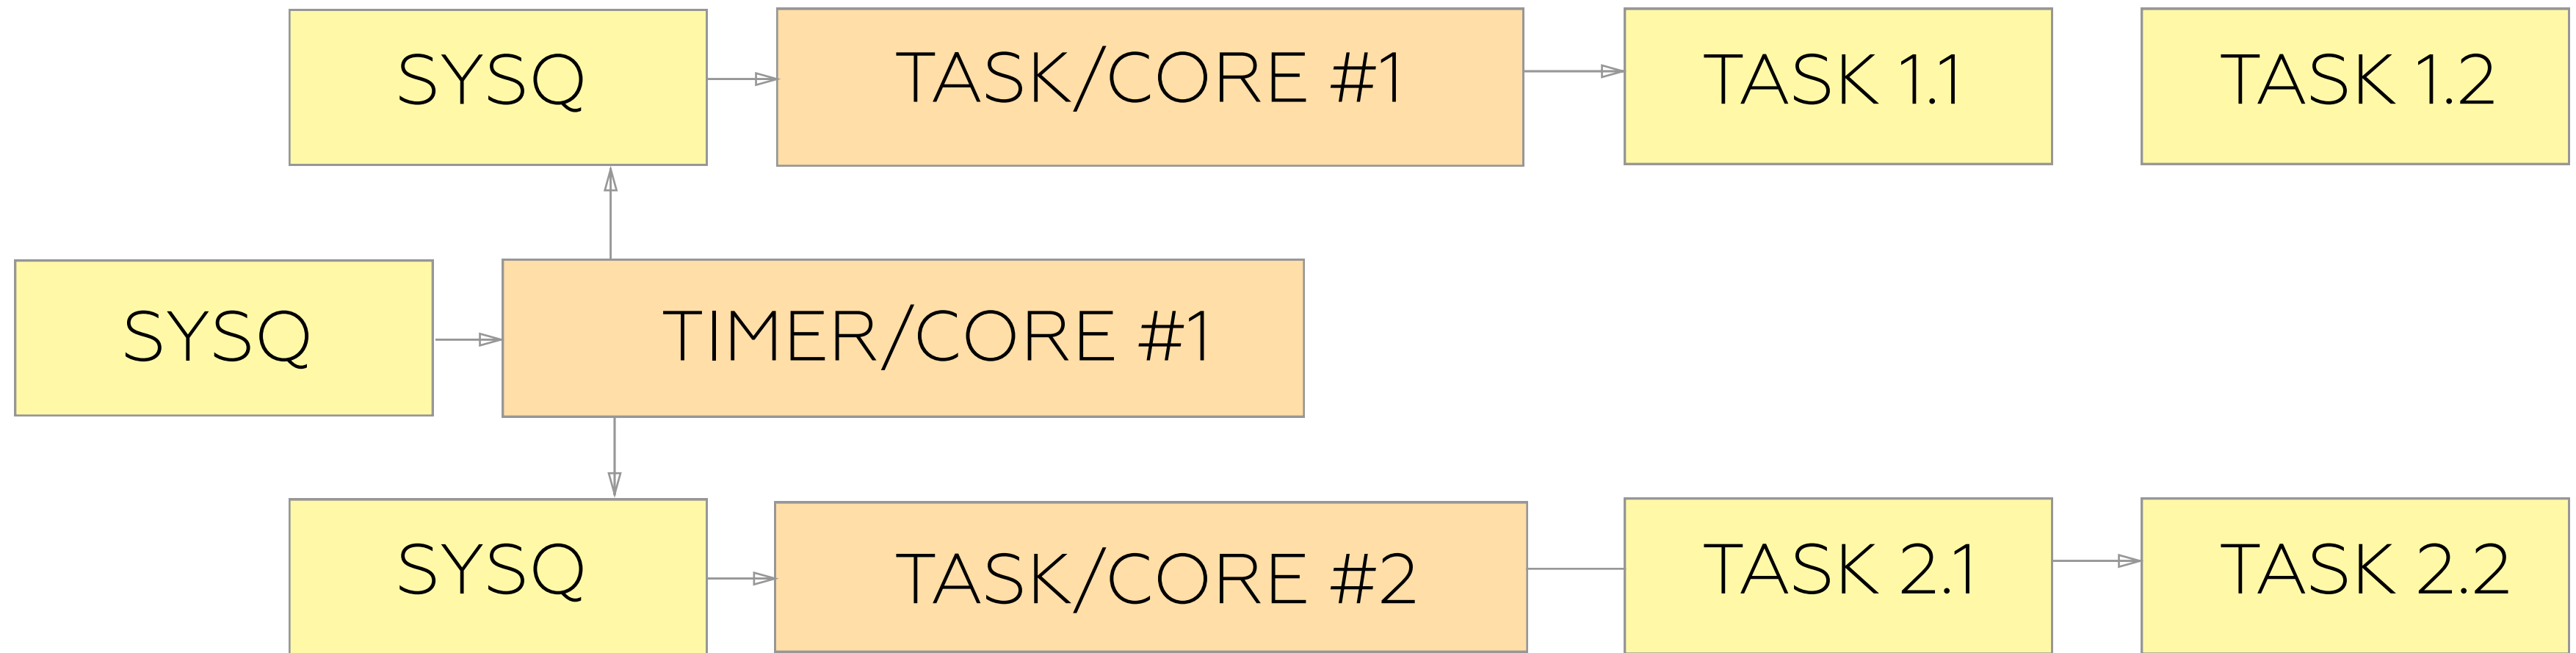
SUBSCRIBER CASE

Multicursor Implementation of SUB (MPSC)
for InterCore Queue Migrations and Cache
Locality



TIMERS

Scheduler Reactors can communicate through InterCore transport for Timers.



Timer uses Linear Firing Round Robin.

Tasks

TASK

STATE VEC

FSM

DATA

CODE

Cursors/Counters

CUR #1 R/W

0—0xFFFF

CUR #2 R

0xFFFF—0xFFFF0000

CUR #3 W

0xFFFF0000—0xFFFFFFFF

CNT #1

00120090912090

Capacity: 239


Time: 20

Workload: 48

Total: 400

Avg Task Consumption
Accumulated in the Task Stream

$$\Sigma \text{Tasks} * \text{Polls} * \text{AvgTime} = \text{Capacity}$$



prios: [10,6,3,1]

ITERATORS

```
+/{x*y}[(1;3;4;5;6);
          (2;6;2;1;3)]
```

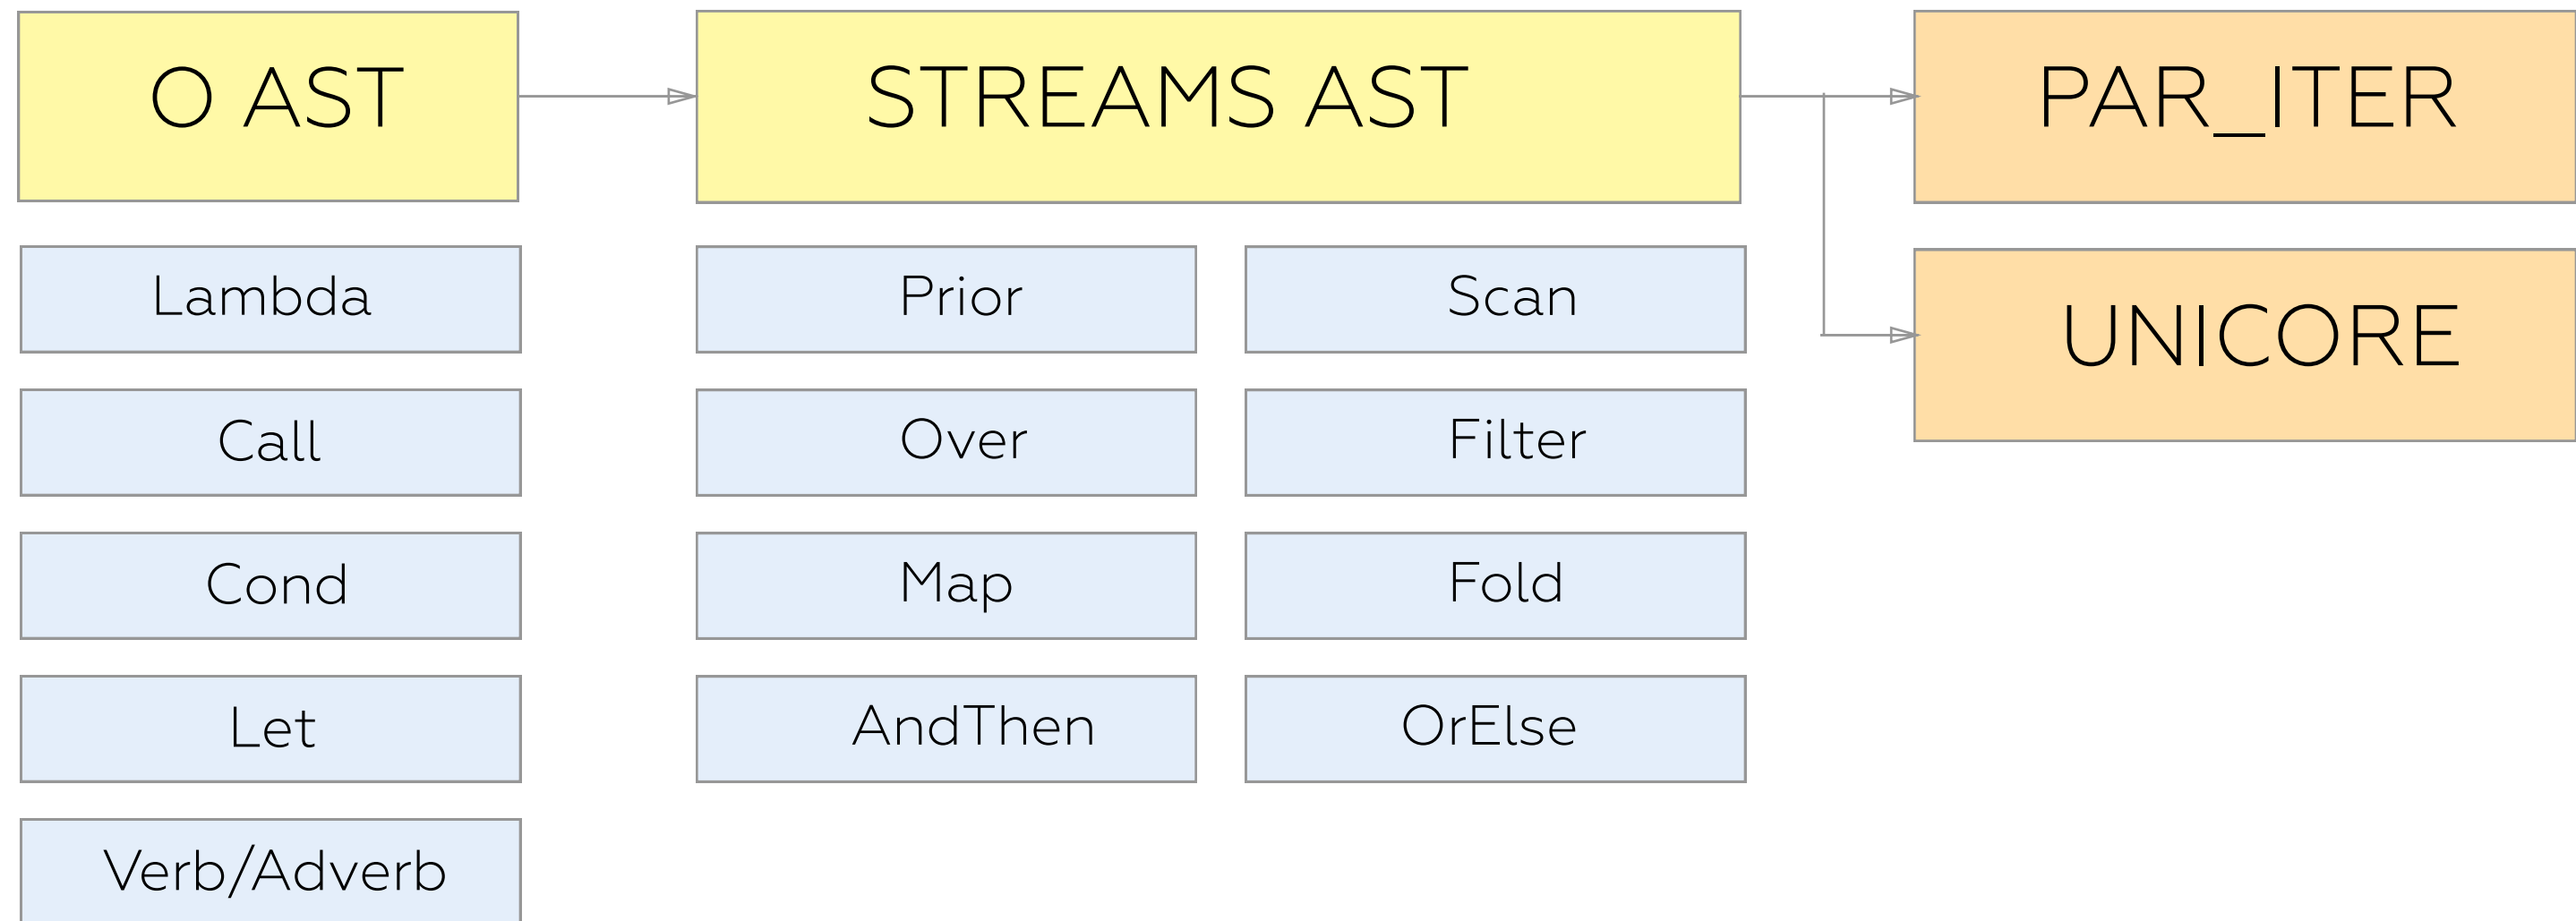
```
+/{x*y}[vec1;vec2]
```

```
vec1.iter().zip(vec2
.iter()).map(|(i, j)|
i * j).sum()
```

```
movdqu    16(%rdx,%rax,4), %xmm2
movdqu    16(%rdi,%rax,4), %xmm3
pshufd    $245, %xmm2, %xmm4
pmuludq   %xmm3, %xmm2
pshufd    $232, %xmm2, %xmm2
pshufd    $245, %xmm3, %xmm3
pmuludq   %xmm4, %xmm3
pshufd    $232, %xmm3, %xmm3
punpckldq %xmm3, %xmm2
paddb     %xmm2, %xmm1
movdqu    (%rdx, %rax,4), %xmm2
movdqu    (%rdi, %rax,4), %xmm3
pshufd    $245, %xmm2, %xmm4
pmuludq   %xmm3, %xmm2
pshufd    $232, %xmm2, %xmm2
pshufd    $245, %xmm3, %xmm3
pmuludq   %xmm4, %xmm3
pshufd    $232, %xmm3, %xmm3
punpckldq %xmm3, %xmm2
paddb     %xmm2, %xmm0
```

INTERPRETER

Unified Pipeline of Language and Streams
Interpretation for Unicore and Multicore



The motivation is to keep LLVM vectorizer continuous happy

```
Console is listening...
>
ring[reader;mem[0;16]];
ring[writer;mem[0;16]];
cursor[1;writer;1];
split[1;2;50];
split[2;3;50];
split[1;4;50];
cursor[5;reader;1];
split[5;6;50];
split[5;7;overlapped];
reactor[aux;0;mod[console;network]];
reactor[timercore;1;mod[timer]];
reactor[core1;2;mod[task]];
reactor[core2;3;mod[task]];
spawn[1;80;AAPL;trader1;core1];
spawn[2;80;EEM-SPY-GDX;trader1;core1];
spawn[3;20;AMI;trader1;core1];
spawn[5;80;GOOG;trader2;core2];
spawn[4;80;FB-NFLX-AMZN;trader2;core2];
timer[timer1;core1;SPY;rule1;t1;notify];
list[reactors];
list[rings];
list[cursors;writer];
list[core1];
list[timercore];
send[1;message1];
send[1;message2];
dump[1;mem[0;100]];
show[recv;1];
```

io

seq

ring

register

spawn

join

send

cursor

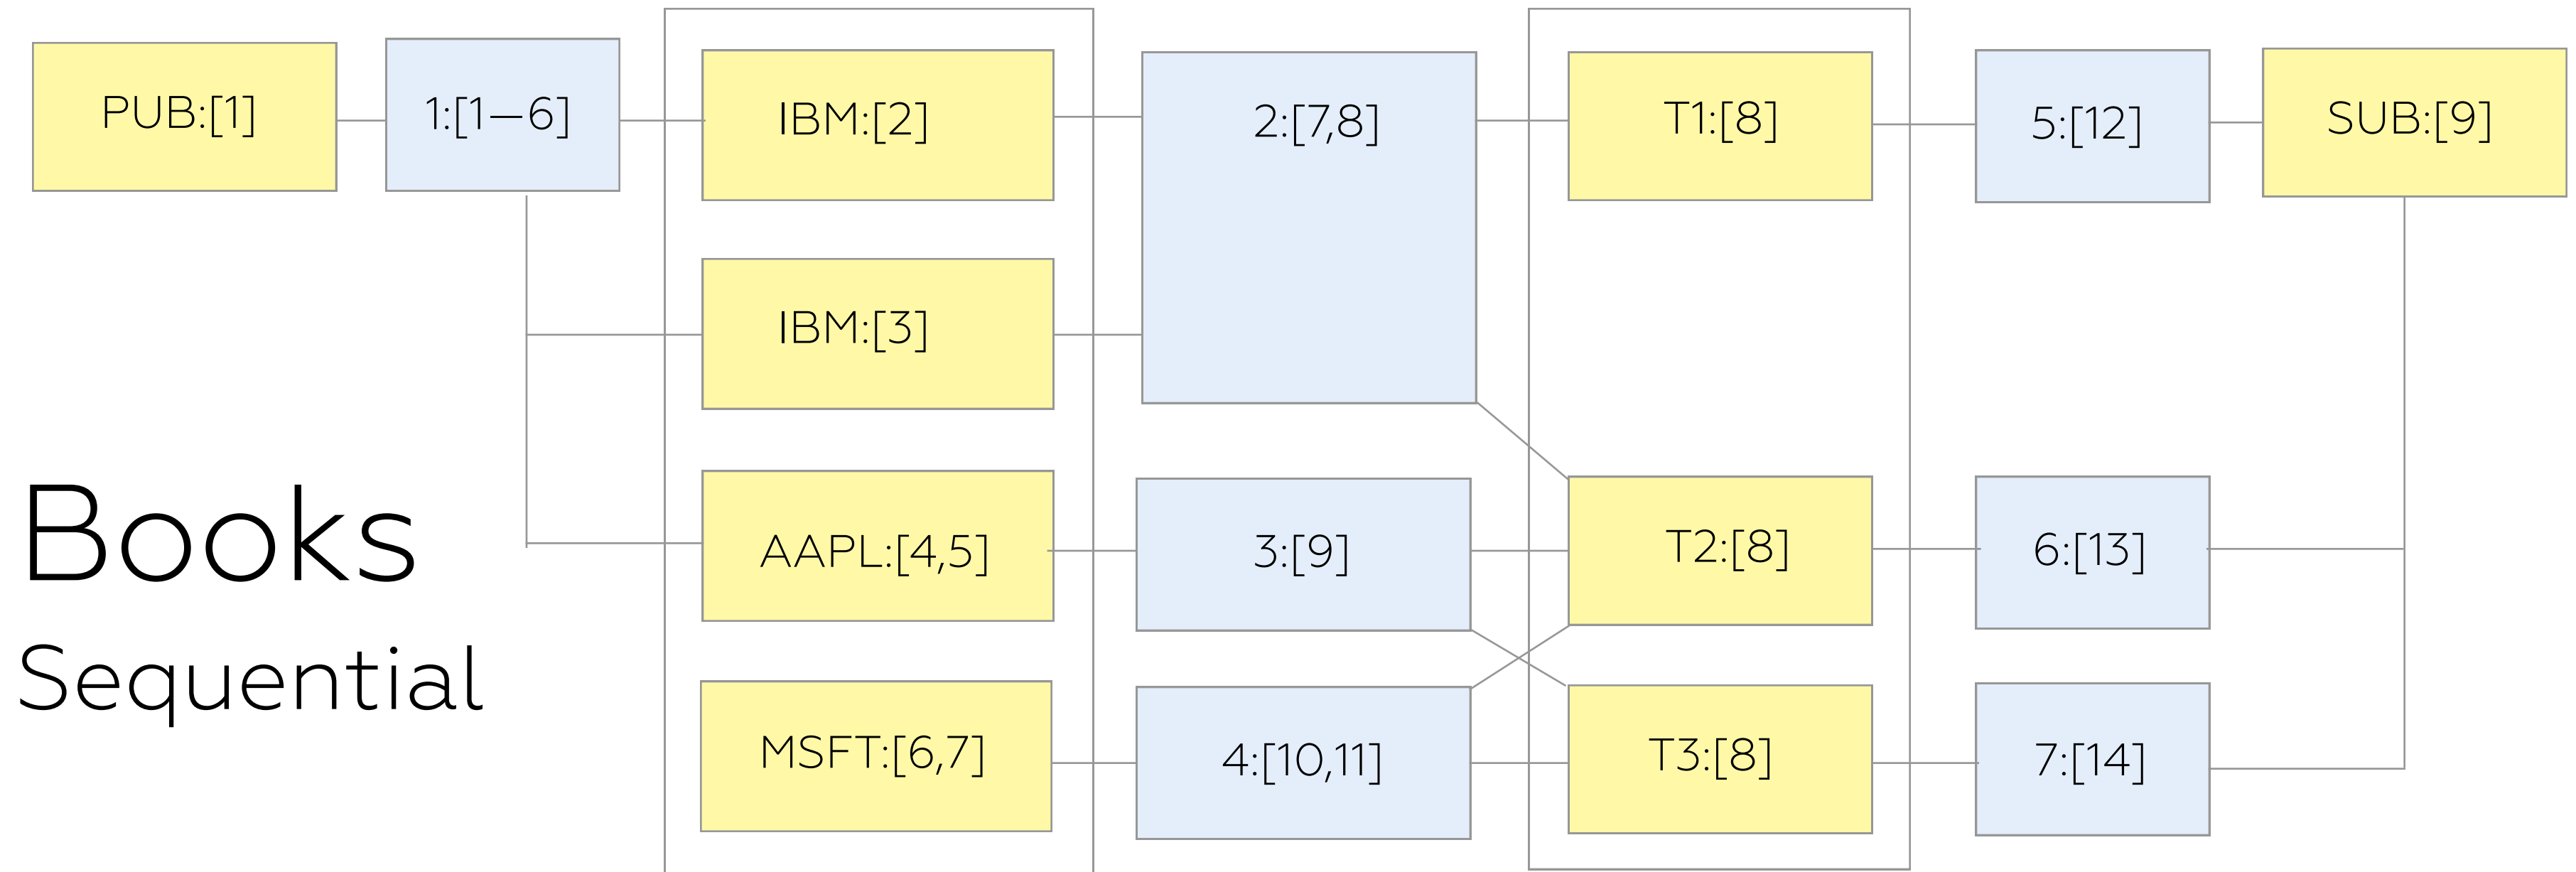
split

sync

reactor

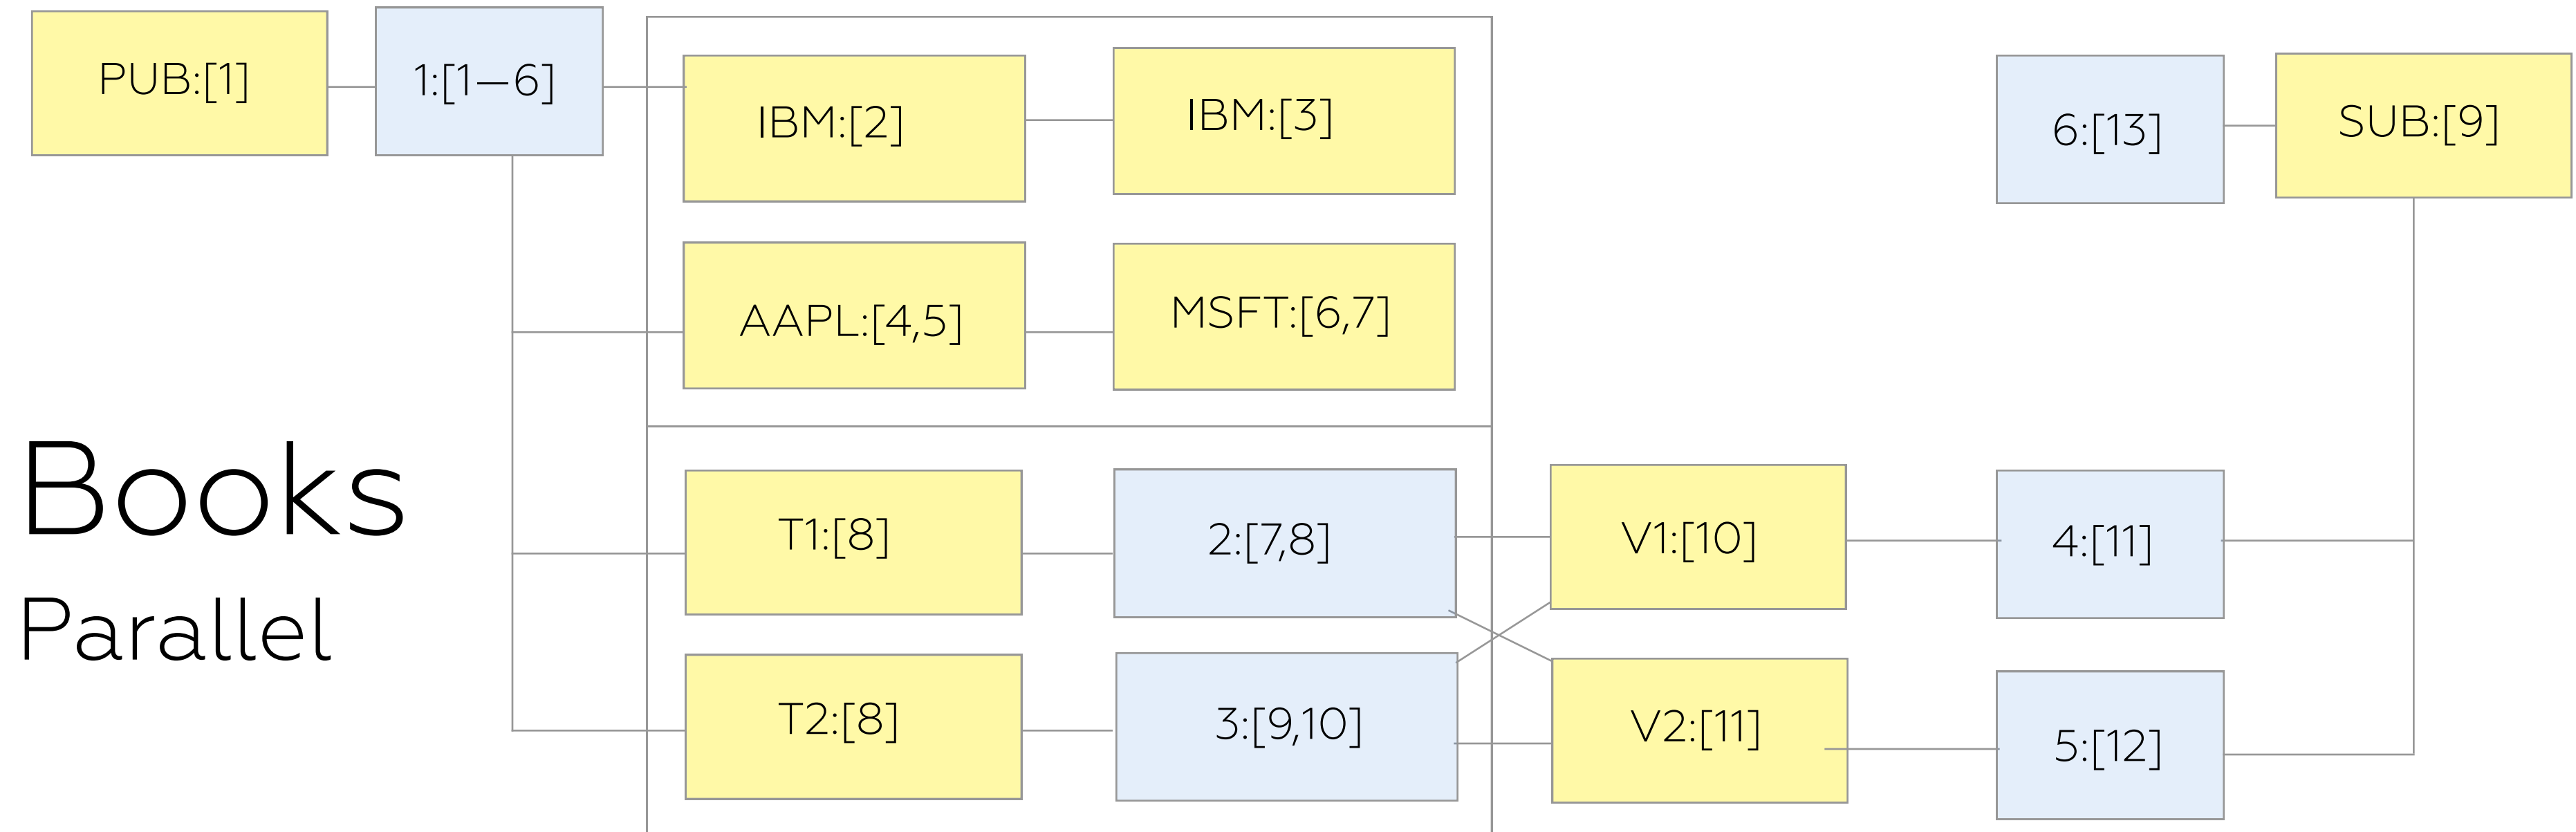
timer

12x CPU Cores: In: [1] Order Books: [2,3,4,5,6,7] Traders: [8] Out: [9]



8x32K MEM Regions: Input Queue: [1] Reducing Queues: [2,3,4,5,6,7]

12x CPU Cores: In: [1] Order Books: [2,3,4,5,6,7] Traders: [8] Out: [9]
Venues: [10,11]



8x32K MEM Regions: Input Queue: [1] Reducing Queues: [2,3,4,5,6]

Book: AAPL

id	side	time	vol	price	venue
====	====	=====	=====	=====	=====
3	ASK	09:05:01:123871012	200	20.30	1
1	ASK	09:01:12:192090139	100	20.30	2
2	ASK	09:03:25:716945237	100	20.25	1
5	BID	09:08:42:134673465	200	20.20	1
4	BID	09:06:11:784316783	100	20.15	1
6	BID	09:09:37:834852874	200	20.15	2