



The Kernel

The Kernel

Scheduler: Round-Robin, Priority Queues, Tree Flavours

Scheduler Actors: Features, Timers, Async I/O

Streams Backends: Zero-copy, Message Passing

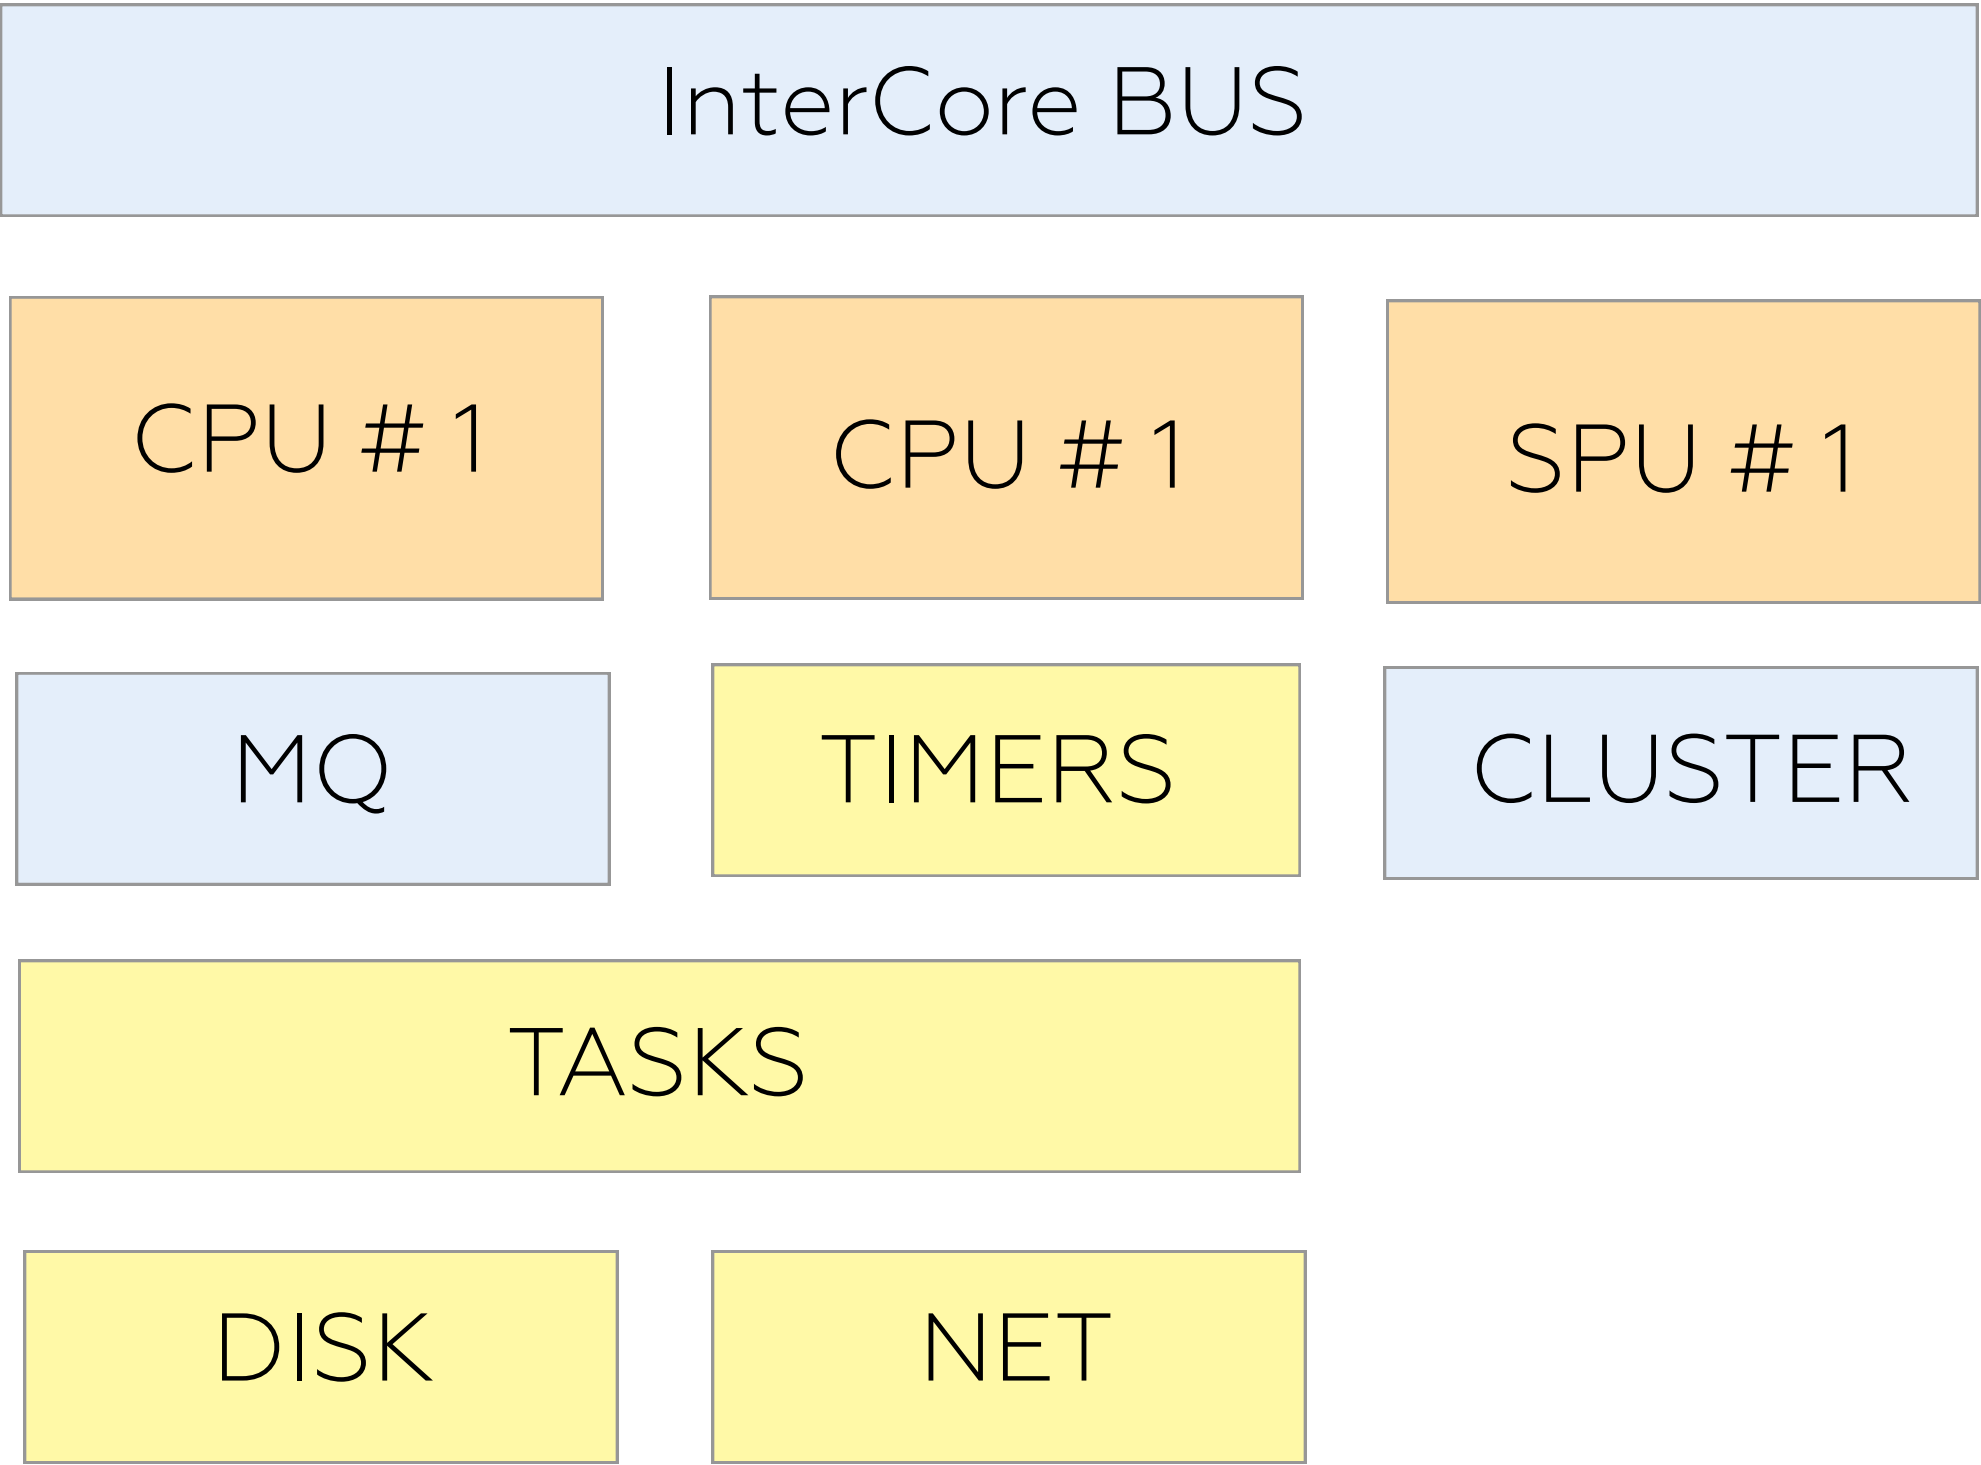
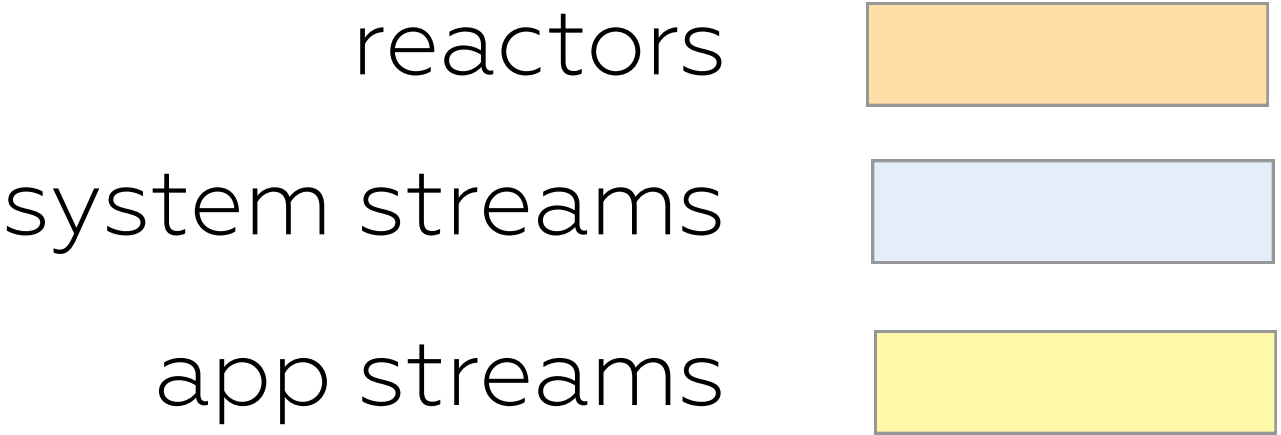
Linear Backends: Async I/O Disk Streams, Network Streams

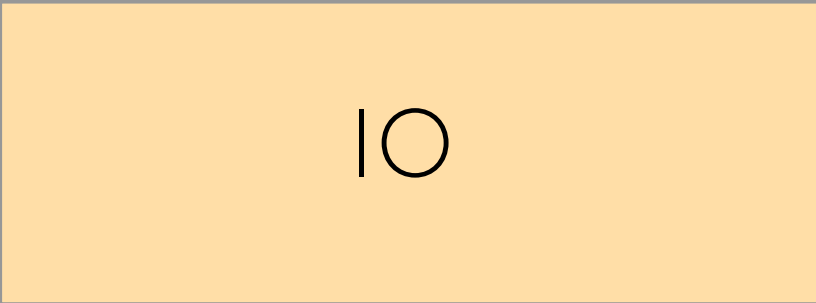
Indexed Backends: Timers, Actors

Backpressured Message Bus/Buffers: Arc/Vec prealloc


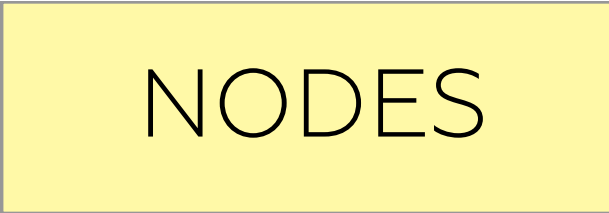
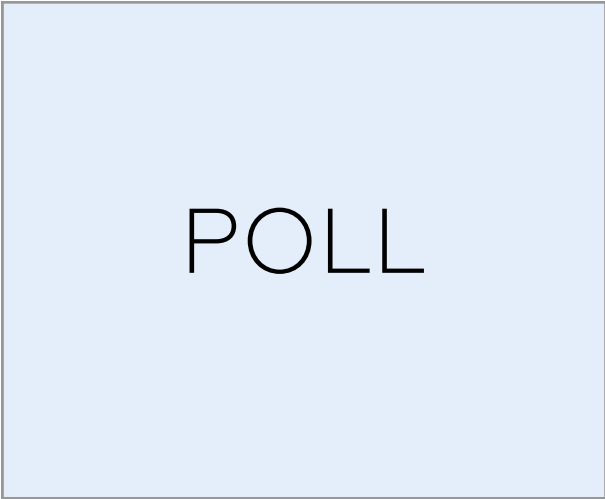
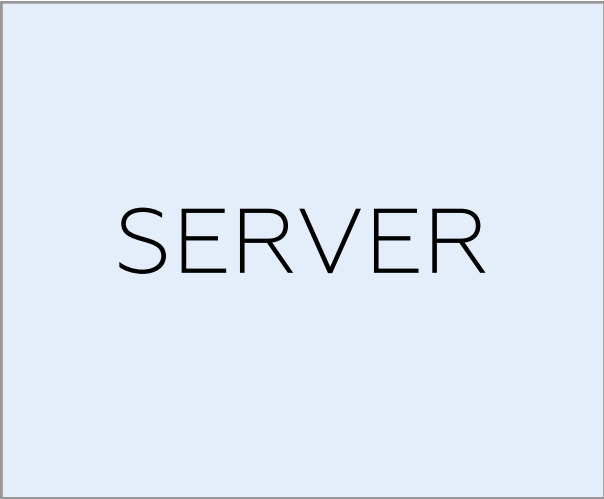
Class: Low Latency, Real Time

Linear: MQ, EXT, DISK, NET
Trees: TIMERS
Priority Queues: TASKS, IRQ





MIO compatible polling loop based on Readiness Queue



OS: EPOLL WAIT

MQ

Queue Types

SPSC/LINK

4-10ns Lowest Latency Possible

MPSC/SUB

10-40ns Reducer or Subscribe Polling

SPMC/PUB

10-40ns Publisher Multicursor

FAST DELIVERY CASE

Single Threaded Task Configuration
to be compared as reference

I/O TASK



CPU TASK



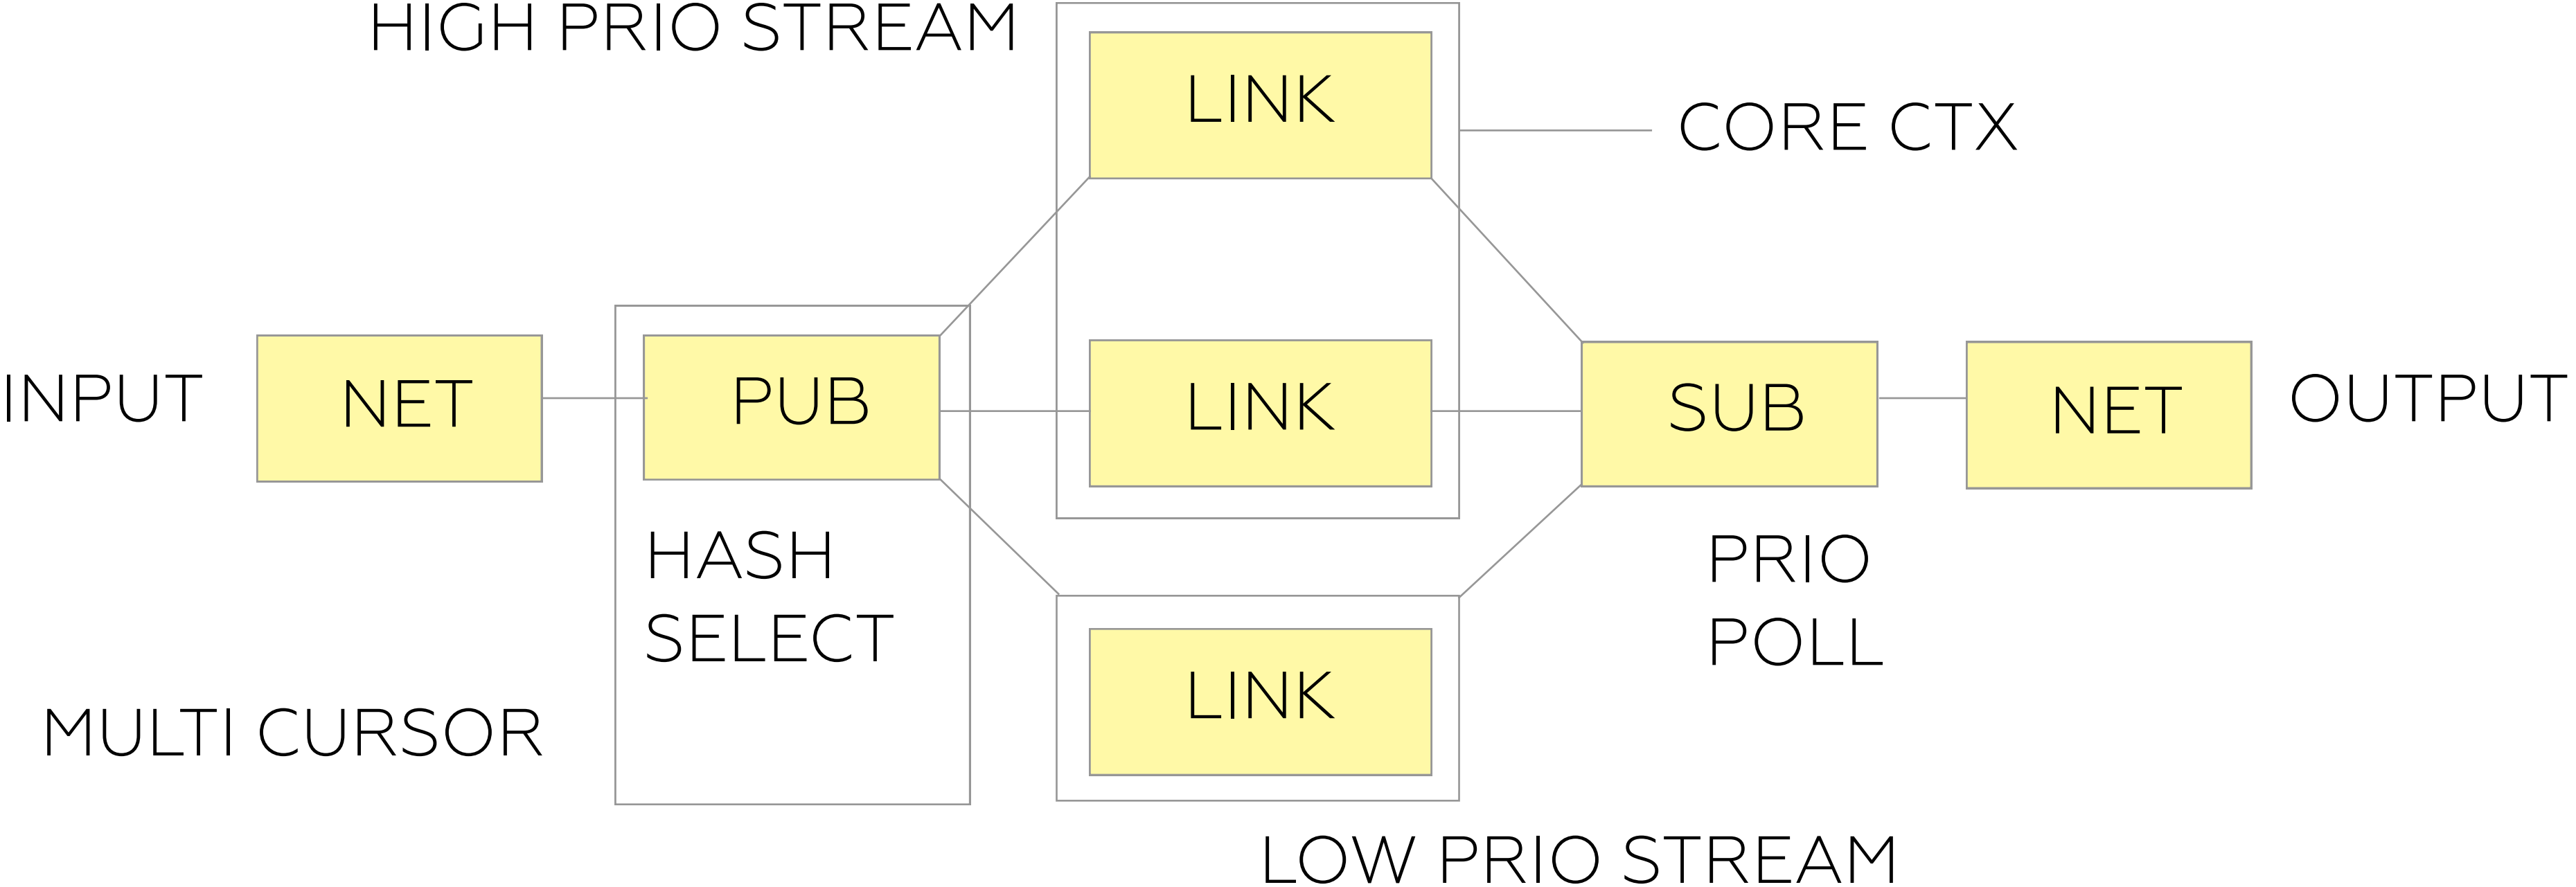
I/O TASK



You can use inplace message modifying and reduce copies to unpack and pack.

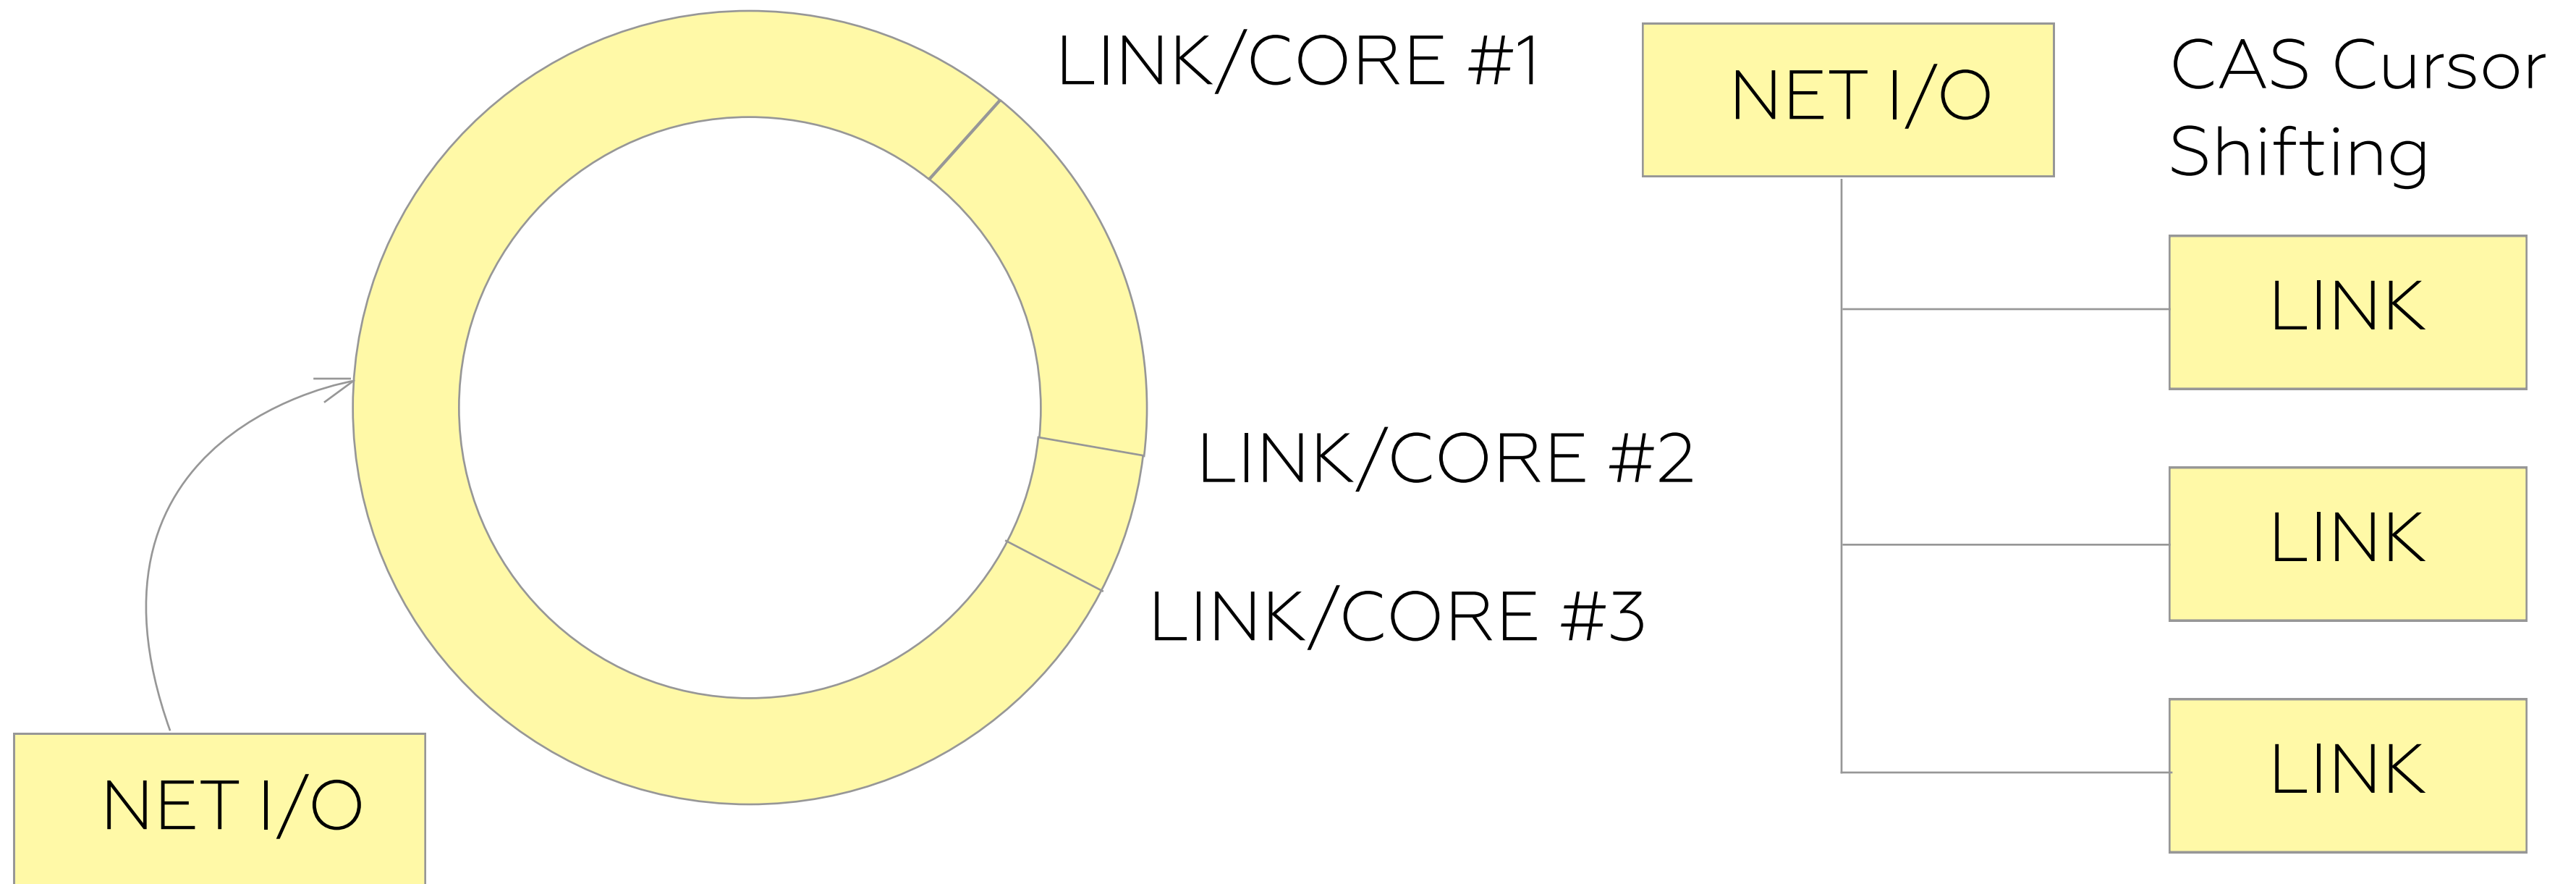
LOAD BALANCING CASE

Load Balancing
of Priority Streams per Core
Buckets



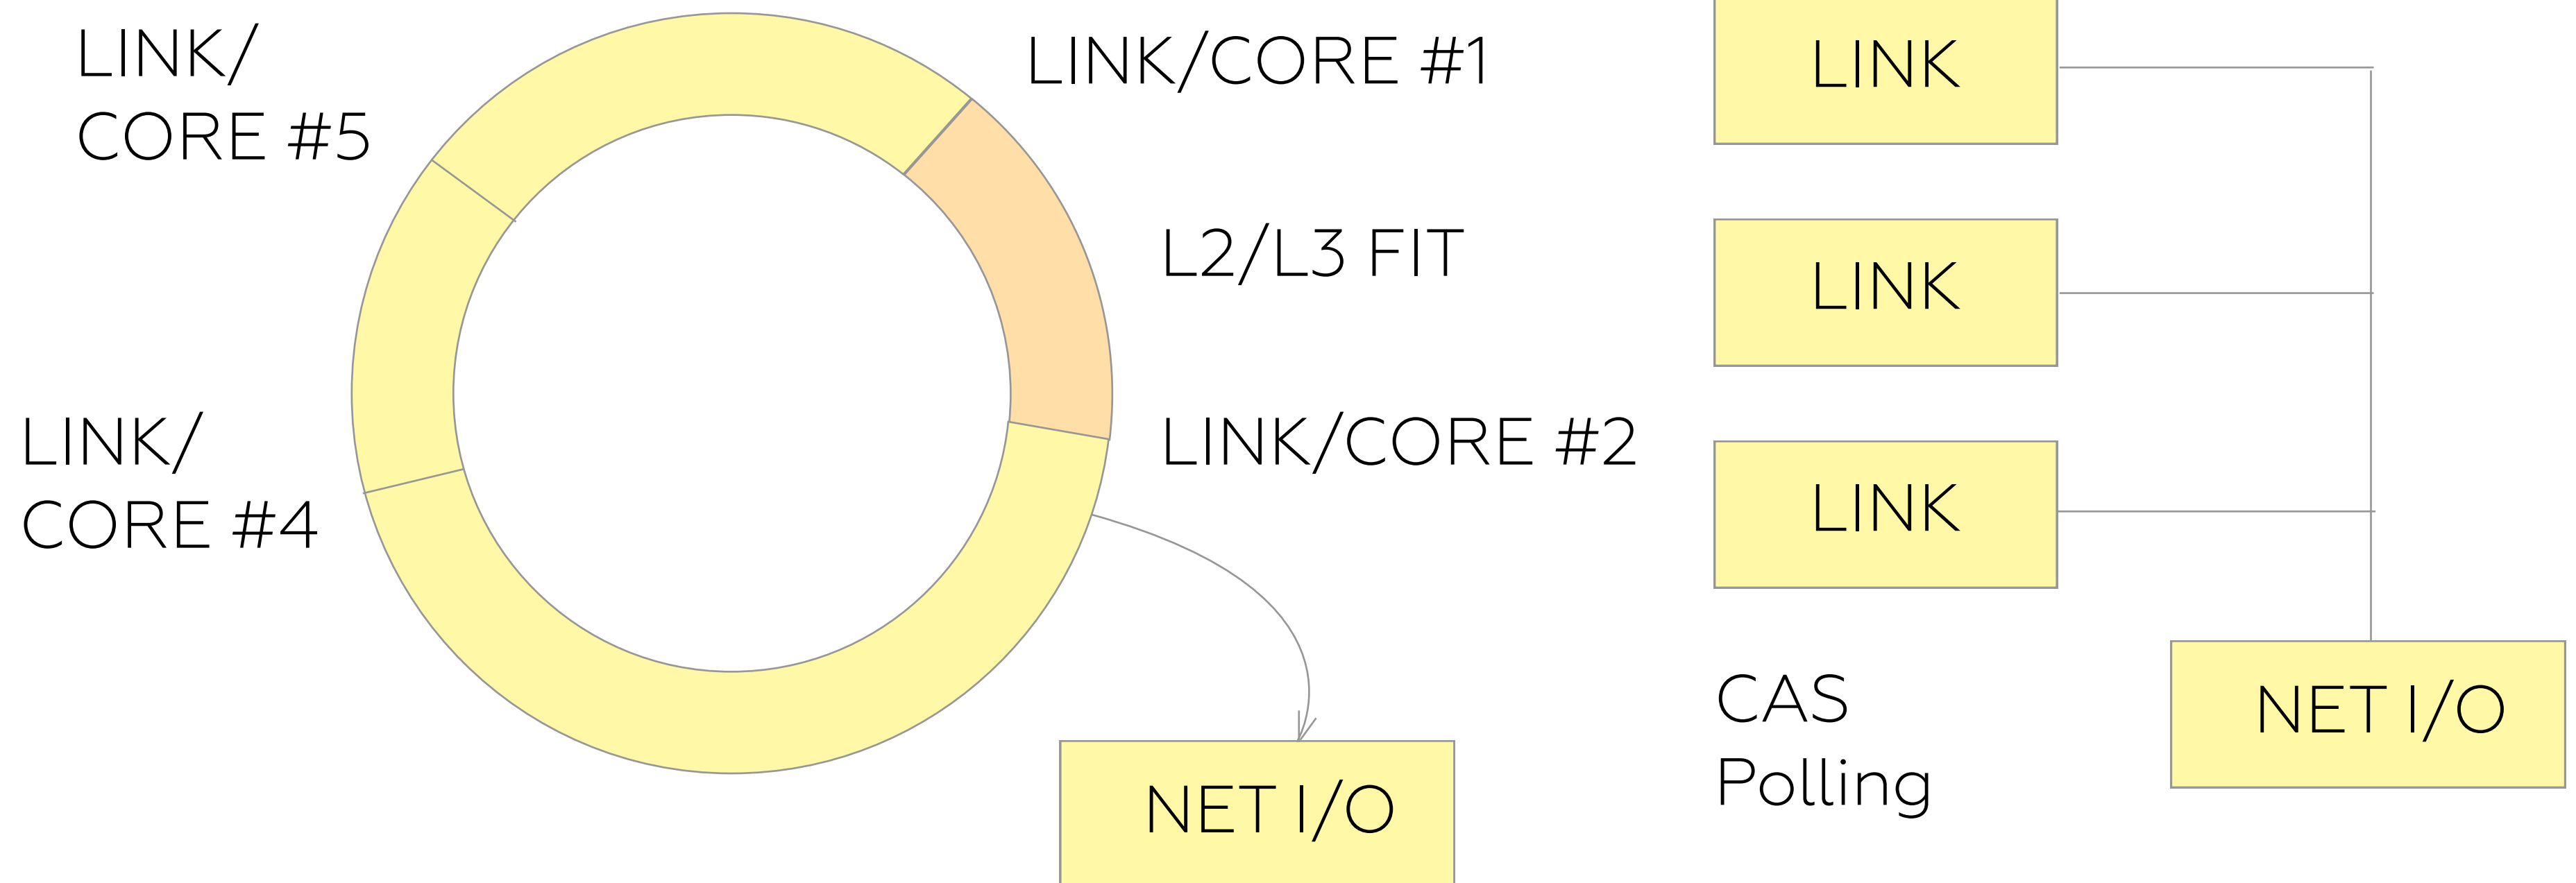
PUBLISHER CASE

PUB Implementation for Zero-Copy
Multiple Consumer Publishing (SPMC)



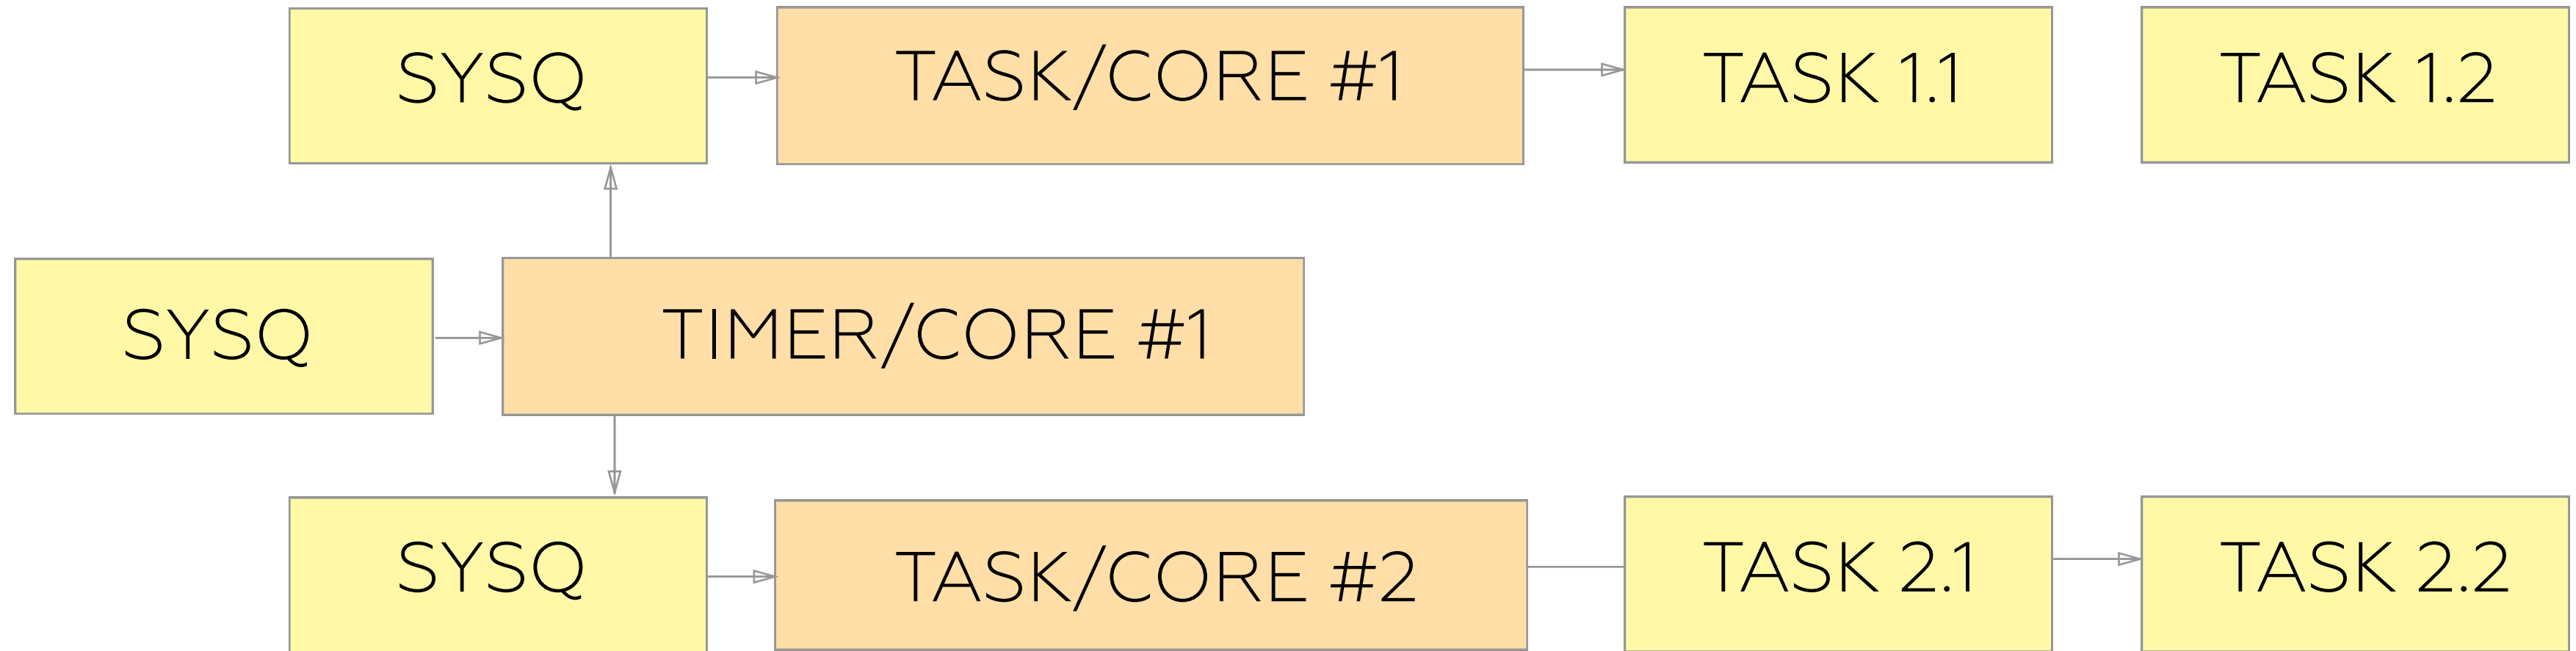
SUBSCRIBER CASE

Multicursor Implementation of SUB (MPSC)
for InterCore Queue Migrations and Cache
Locality



TIMERS

Scheduler Reactors can communicate through InterCore transport for Timers.



Timer uses Linear Firing Round Robin.

Tasks

Cursors/Counters

TASK

STATE VEC

FSM

DATA

CODE

CUR #1 R/W

CUR #2 R

CUR #3 W

CNT #1

0—0xFFFF

0xFFFF—0xFFFF0000

0xFFFF0000—0xFFFFFFFF

00120090912090

ITERATORS

```
+/{x*y}[(1;3;4;5;6);  
         (2;6;2;1;3)]
```

```
+/{x*y}[vec1;vec2]
```

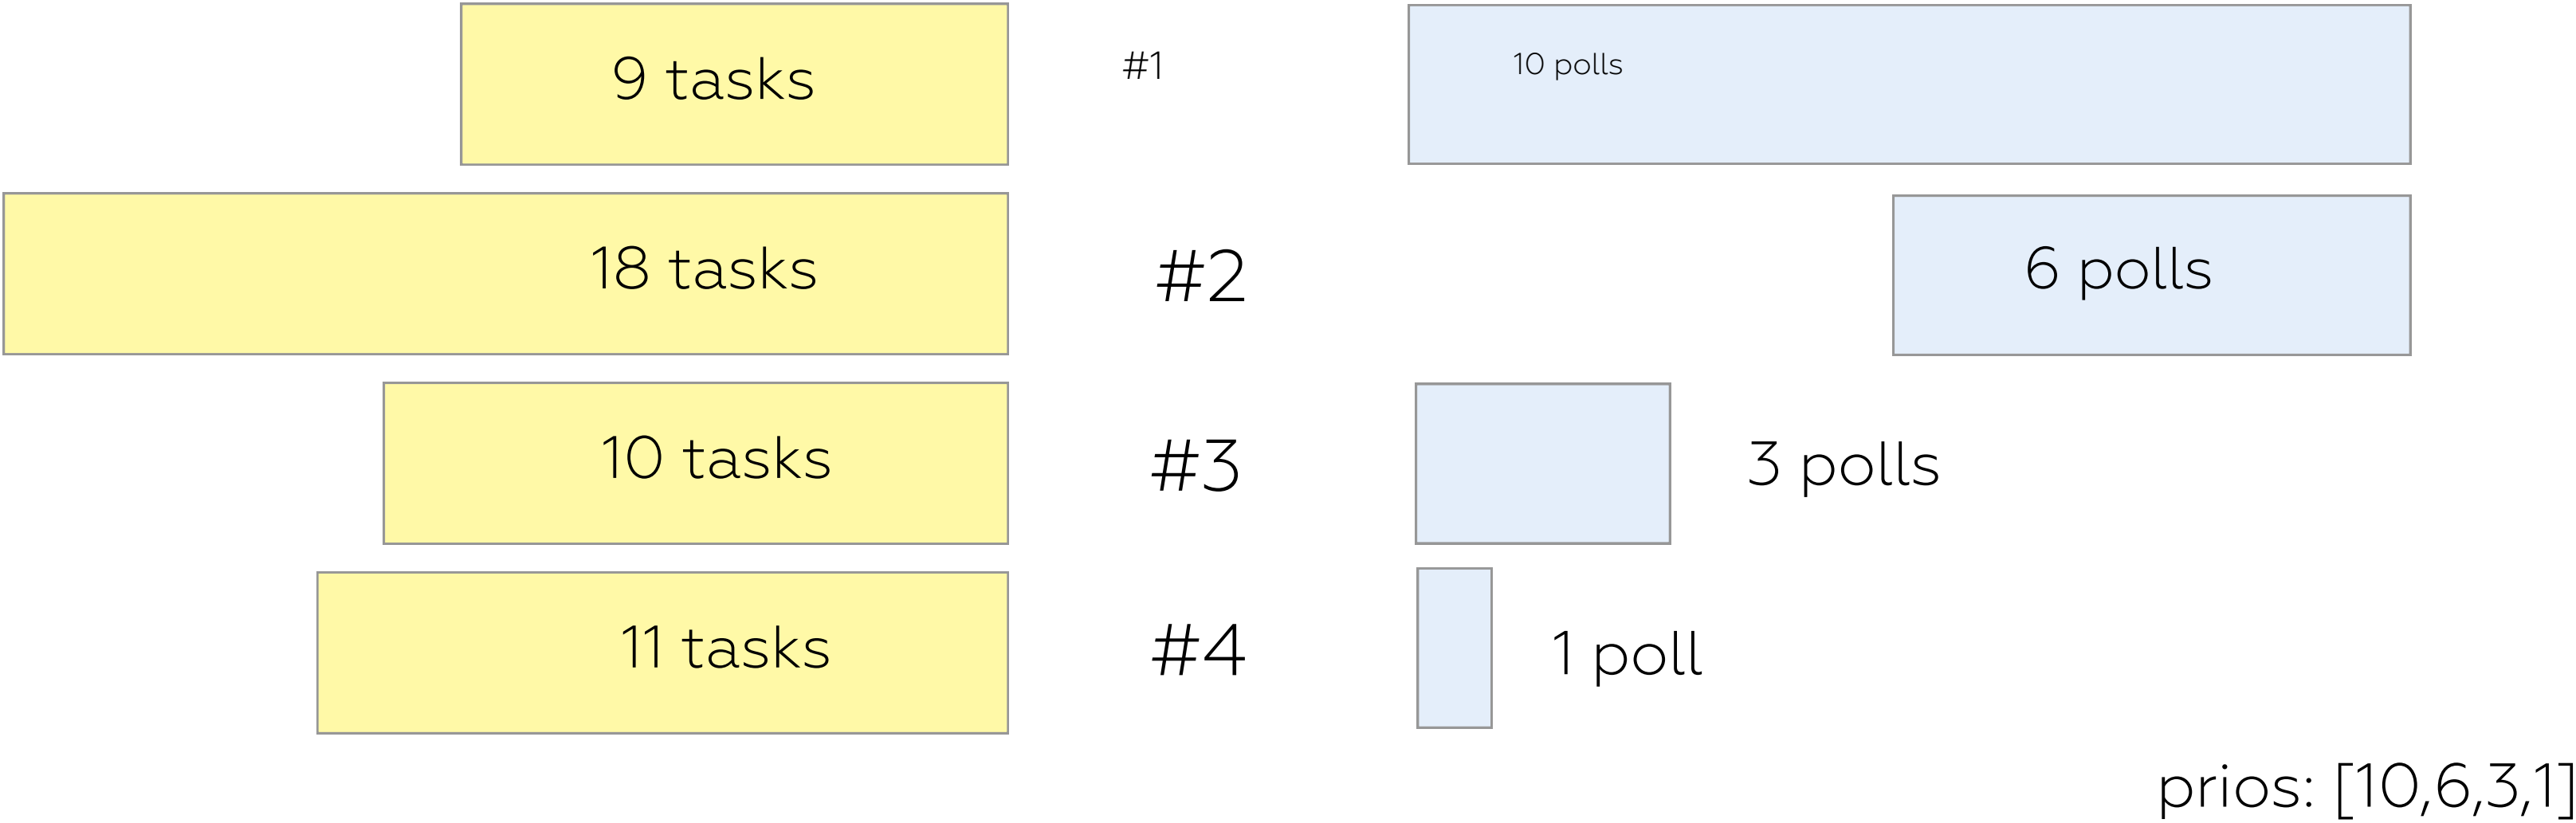
```
vec1.iter().zip(vec2  
  .iter()).map(|(i, j)|  
  i * j).sum()
```

```
.LBB0_7: testq %r8, %r8  
        je .LBB0_9  
        movdqu 16(%rdx,%rax,4), %xmm2  
        movdqu 16(%rdi,%rax,4), %xmm3  
        pshufd $245, %xmm2, %xmm4  
        pmuludq %xmm3, %xmm2  
        pshufd $232, %xmm2, %xmm2  
        pshufd $245, %xmm3, %xmm3  
        pmuludq %xmm4, %xmm3  
        pshufd $232, %xmm3, %xmm3  
        punpckldq %xmm3, %xmm2  
        padd %xmm2, %xmm1  
        movdqu (%rdx,%rax,4), %xmm2  
        movdqu (%rdi,%rax,4), %xmm3  
        pshufd $245, %xmm2, %xmm4  
        pmuludq %xmm3, %xmm2  
        pshufd $232, %xmm2, %xmm2  
        pshufd $245, %xmm3, %xmm3  
        pmuludq %xmm4, %xmm3  
        pshufd $232, %xmm3, %xmm3  
        punpckldq %xmm3, %xmm2  
        padd %xmm2, %xmm0
```

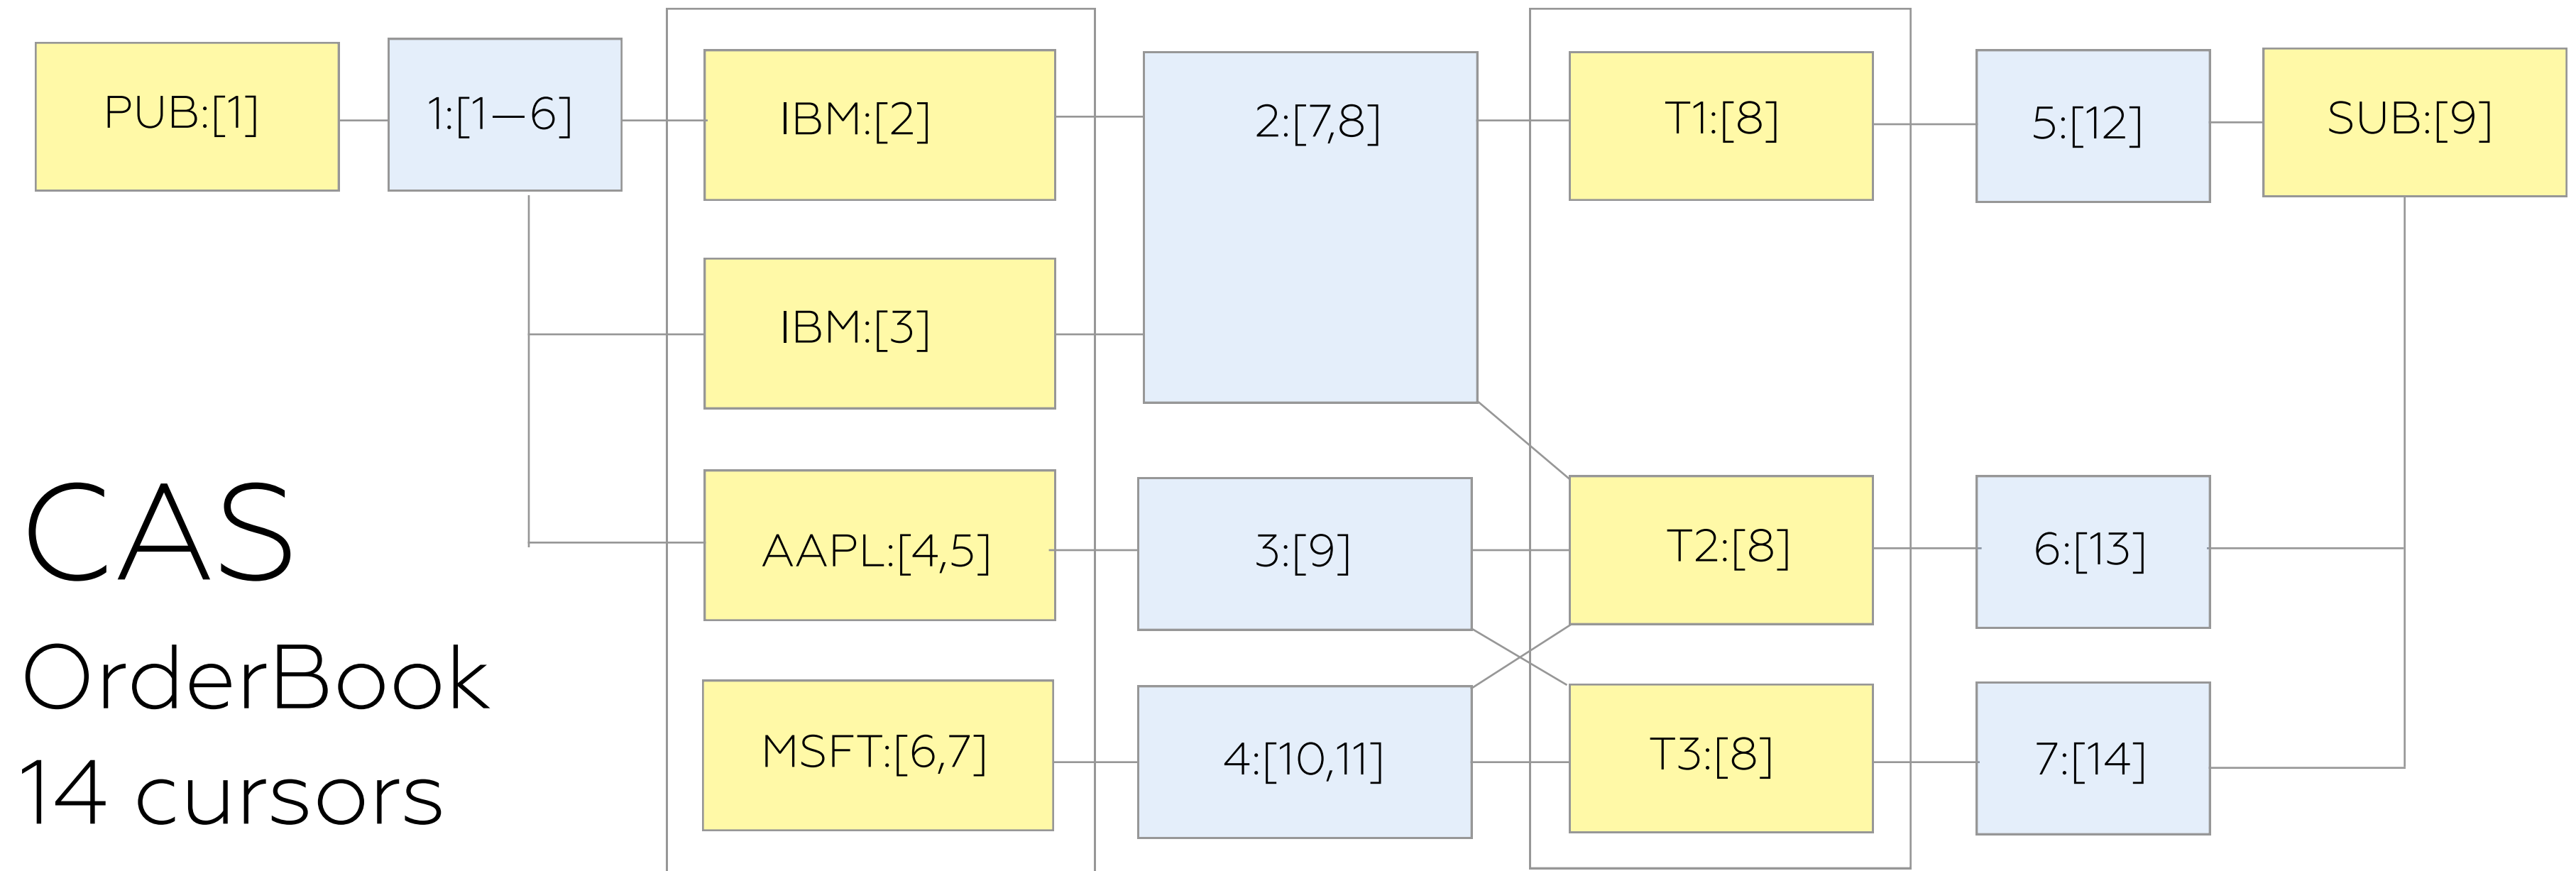
Capacity: 239 Time: 20
Workload: 48 Total: 400

Avg Task Consumption
Accumulated in the Task Stream

$\Sigma \text{ AvgTime} * \text{Tasks} * \text{Polls} = \text{Capacity}$



12x CPU Cores: In: [1] Order Books: [2,3,4,5,6,7] Traders: [8] Out: [9]



CAS
OrderBook
14 cursors

8x32K MEM Regions: Input Queue: [1] Reducing Queues: [2,3,4,5,6,7]

```
Console is listening...
>
ring[reader;mem[0;16]];
ring[writer;mem[0;16]];
cursor[1;writer;1];
split[1;2;50];
split[2;3;50];
split[1;4;50];
cursor[5;reader;1];
split[5;6;50];
split[5;7;overlapped];
reactor[aux;0;mod[console;network]];
reactor[timercore;1;mod[timer]];
reactor[core1;2;mod[task]];
reactor[core2;3;mod[task]];
spawn[1;80;AAPL;trader1;core1];
spawn[2;80;EEM-SPY-GDX;trader1;core1];
spawn[3;20;AMI;trader1;core1];
spawn[5;80;GOOG;trader2;core2];
spawn[4;80;FB-NFLX-AMZN;trader2;core2];
timer[timer1;core1;SPY;rule1;t1;notify];
list[reactors];
list[rings];
list[cursors;writer];
list[core1];
list[timercore];
send[1;message1];
send[1;message2];
dump[1;mem[0;100]];
show[recv;1];
```

io	seq	ring
register	spawn	join
send	cursor	split
sync	reactor	timer